

# Practice Project - Linear Regression based : Using dataset (Advertising.csv)

## Table of Contents

Problem statement

1. Introduction

2. Linear Regression Steps

2.1 Importing Our Libraries

2.2 Reading the data

2.3 Assumptions Check

2.4 Preprocessing and Understanding the data

2.5 Splitting the data into train and test

2.6 Running your First Machine Learning Algorithm

2.7 Evaluation

2.7.1 RMSE

2.7.2 R Squared

2.8 Finalizing Stuff

## 1. Introduction

- Regression is a statistical technique which helps you to measure the relationship between the independent variables and dependent variables
- It helps you to understand one unit change in the independent variables is going to cause how many units change in the dependent variable
- Dependent or predicted variable is represented as 'y'

## 2. Linear Regression Steps:

### 2.1 Importing Our Libraries

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

### 2.2 Reading the data

Here two important things

1. index\_col is by default none and header is 0

1. Make use of proper function depending on the extension of the file

```
In [ ]: data = pd.read_csv(r'Advertising.csv', index_col = 0)
```

```
In [ ]: data.head()
```

```
Out [ ]:
```

	TV	Radio	Newspaper	Sales
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
3	17.2	45.9	69.3	9.3
4	151.5	41.3	58.5	18.5
5	180.8	10.8	58.4	12.9

## 2.3 Assumptions Check

Assumption 1: There should be no outliers in the data

```
In [ ]: data.describe()
```

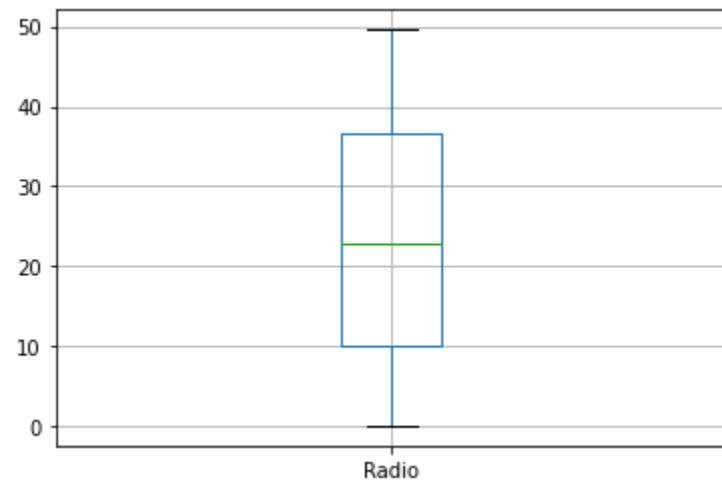

```
Out [ ]:
```

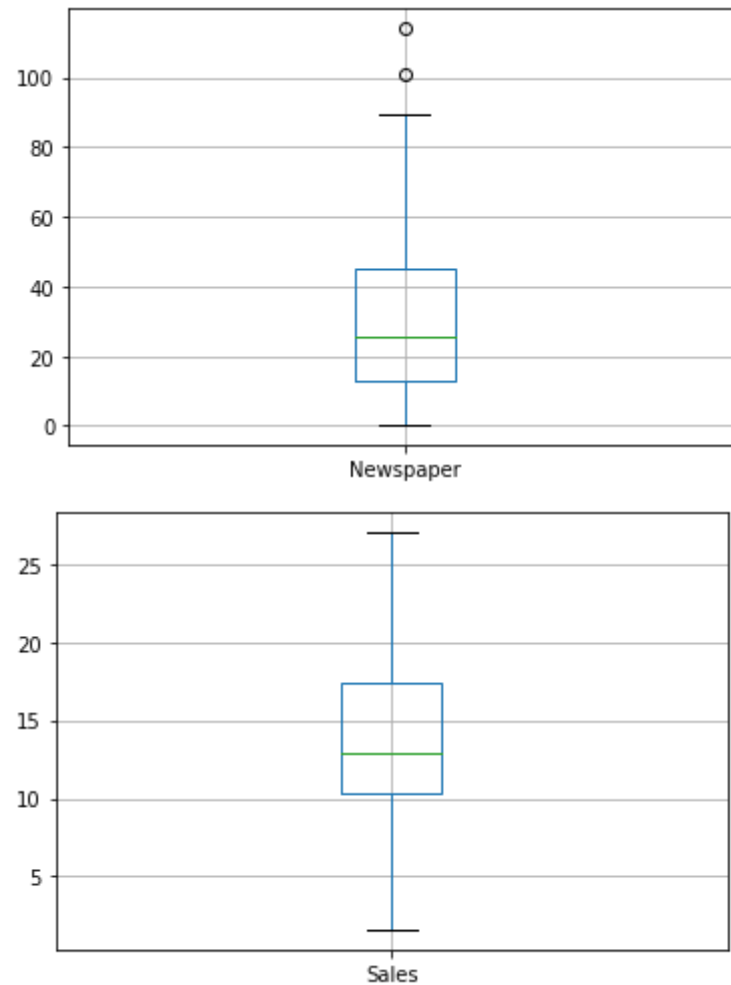
	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	14.022500
std	85.854236	14.846809	21.778621	5.217457
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	10.375000
50%	149.750000	22.900000	25.750000	12.900000

	TV	Radio	Newspaper	Sales
<b>75%</b>	218.825000	36.525000	45.100000	17.400000
<b>max</b>	296.400000	49.600000	114.000000	27.000000

In [ ]:

```
for i in data.columns:  
    data.boxplot(column=i)  
    plt.show()
```






Assumption 2: Assumption of Linearity : Every independent variable should have a linear relationship with the dependent variable


```
In [ ]: sns.pairplot(data,x_vars=['TV','Radio','Newspaper'],  
                  y_vars="Sales",kind='reg')
```

```
Out[ ]: <seaborn.axisgrid.PairGrid at 0x251f619a040>
```



Assumption 3: Assumption of no multicollinearity means the independent variable should not be correlated with each other

```
In [ ]: corr = data.corr()
sns.heatmap(corr, annot=True, square=True)
plt.xticks(rotation=0)
sns.set(rc = {'figure.figsize':(15,8)})
plt.show()
```



## 2.4 Preprocessing and Understanding the data

```
In [ ]: data.shape
```

Out[ ]: (200, 4)

In [ ]: `data.isnull()`

Out[ ]:

	TV	Radio	Newspaper	Sales
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
5	False	False	False	False
...	...	...	...	...
196	False	False	False	False
197	False	False	False	False
198	False	False	False	False
199	False	False	False	False
200	False	False	False	False

200 rows × 4 columns

In [ ]: `data.isnull().sum()`

Out[ ]: TV 0  
Radio 0  
Newspaper 0  
Sales 0  
dtype: int64

So as there are no null values in the data we don't need any sort of missing values handling but let's consider an example where we have missing values and we treat it using mean, median, and mode or we can delete the column itself.

## 2.5 Splitting the data into train and test

```
In [ ]: x = data[['TV', 'Radio', 'Newspaper']]
        y = data[['Sales']]
```

```
In [ ]: import sklearn
```

```
In [ ]: from sklearn.model_selection import train_test_split
        x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=10)
```

```
In [ ]: print(x_train.shape)
        print(x_test.shape)
```

```
(160, 3)
(40, 3)
```

```
In [ ]: print(y_train.shape)
        print(y_test.shape)
```

```
(160, 1)
(40, 1)
```

## 2.6 Running your First Machine Learning Algorithm

```
In [ ]: from sklearn.linear_model import LinearRegression

        lm =LinearRegression()

        lm.fit(x_train,y_train)
```

```
Out[ ]: LinearRegression()
```

```
In [ ]: print(lm.intercept_)
```

```
[3.25409711]
```

```
In [ ]: print(lm.coef_)
```

```
[[ 0.0437726  0.19343299 -0.00222879]]
```

```
In [ ]: y_pred = lm.predict(x_test)
print(y_pred)
```

```
[[18.1625299 ]
 [12.92663232]
 [18.0531098 ]
 [23.64464668]
 [20.70438374]
 [14.28227997]
 [14.94493548]
 [21.38232981]
 [21.17508238]
 [12.73110461]
 [24.00312134]
 [ 7.21544071]
 [12.24762152]
 [19.24345998]
 [19.38241854]
 [13.45643798]
 [19.6247089 ]
 [ 9.2531648 ]
 [21.13268075]
 [20.90762408]
 [15.53485445]
 [10.92529369]
 [22.82955184]
 [15.8122438 ]
 [17.42515749]
 [ 8.16218669]
 [11.89783444]
 [12.70337575]
 [21.74138085]
 [ 7.96215368]
 [12.50099965]
 [20.45535282]
 [ 4.72120047]
 [ 4.72259288]
 [16.75292333]
 [15.75804986]
 [ 6.74415499]
 [17.73477354]
```



```
[ 9.01591827]  
[13.617599   ]]
```

## 2.7 Evaluation

### 2.7.1 RMSE:

- Root Mean Square error is an absolute measure of the goodness for the fit
- It gives an absolute number on how much your predicted results deviate from the actual number
- Low the RMSE better the model

### 2.7.2 R Squared:

- It tells you how well the regression model is predicting as compared to the mean model
- Lies between (0-1)
- If R squared is close to 1 → very good model

```
In [ ]: from sklearn.metrics import r2_score, mean_squared_error  
  
r2 = r2_score(y_test, y_pred)  
  
rmse = np.sqrt(mean_squared_error(y_test, y_pred))  
  
print('R-squared', r2)  
print('RMSE', rmse)
```

```
R-squared 0.8353672324670594  
RMSE 2.58852984462781
```

## 2.8 Finalizing Stuff

```
In [ ]: new_data = pd.DataFrame()  
new_data = (x_test)  
  
new_data['Actual_Sales'] = y_test  
new_data['Predicted_Sales'] = y_pred
```

```
In [ ]: new_data
```

Out[ ]:

	TV	Radio	Newspaper	Actual_Sales	Predicted_Sales
<b>60</b>	210.7	29.5	9.3	18.4	18.162530
<b>6</b>	8.7	48.9	75.0	7.2	12.926632
<b>21</b>	218.4	27.7	53.4	18.0	18.053110
<b>199</b>	283.6	42.0	66.2	25.5	23.644647
<b>53</b>	216.4	41.7	39.6	22.6	20.704384
<b>20</b>	147.3	23.9	19.1	14.6	14.282280
<b>163</b>	188.4	18.1	25.6	14.9	14.944935
<b>56</b>	198.9	49.4	60.0	23.7	21.382330
<b>70</b>	216.8	43.9	27.2	22.3	21.175082
<b>3</b>	17.2	45.9	69.3	9.3	12.731105
<b>99</b>	289.7	42.3	51.2	25.4	24.003121
<b>11</b>	66.1	5.8	24.2	8.6	7.215441
<b>76</b>	16.9	43.7	89.4	8.7	12.247622
<b>143</b>	220.5	33.2	37.9	20.1	19.243460
<b>125</b>	229.5	32.3	74.2	19.7	19.382419
<b>64</b>	102.7	29.6	8.4	14.0	13.456438
<b>110</b>	255.4	26.9	5.5	19.8	19.624709
<b>79</b>	5.4	29.9	9.4	5.3	9.253165
<b>112</b>	241.7	38.0	23.2	21.8	21.132681
<b>186</b>	205.0	45.1	19.6	22.6	20.907624
<b>155</b>	187.8	21.1	9.5	15.6	15.534854
<b>131</b>	0.7	39.6	8.7	1.6	10.925294
<b>62</b>	261.3	42.7	54.7	24.2	22.829552
<b>88</b>	110.7	40.6	63.2	16.0	15.812244

	TV	Radio	Newspaper	Actual_Sales	Predicted_Sales
<b>103</b>	280.2	10.1	21.4	14.8	17.425157
<b>122</b>	18.8	21.7	50.4	7.0	8.162187
<b>137</b>	25.6	39.0	9.3	9.5	11.897834
<b>2</b>	44.5	39.3	45.1	10.4	12.703376
<b>48</b>	239.9	41.5	18.5	23.2	21.741381
<b>173</b>	19.6	20.1	17.0	7.6	7.962154
<b>160</b>	131.7	18.4	34.6	12.9	12.501000
<b>40</b>	228.0	37.7	32.0	21.5	20.455353
<b>77</b>	27.5	1.6	20.7	6.9	4.721200
<b>92</b>	28.6	1.5	33.0	7.3	4.722593
<b>36</b>	290.7	4.1	8.5	12.8	16.752923
<b>179</b>	276.7	2.3	23.7	11.8	15.758050
<b>128</b>	80.2	0.0	9.2	8.8	6.744155
<b>170</b>	284.3	10.6	6.4	15.0	17.734774
<b>47</b>	89.7	9.9	35.7	10.6	9.015918
<b>175</b>	222.4	3.4	13.1	11.5	13.617599

In [ ]:

```
#new_data.to_csv('predictions.csv')
```