

# Practice Project - Dataframe based: Advanced GroupBy

## Table of Contents

- Problem statement
- 1. Importing Libraries
- 2. Python implementation

## Groupby operations

Some imports:

```
In [ ]: %matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
try:
    import seaborn
except ImportError:
    pass

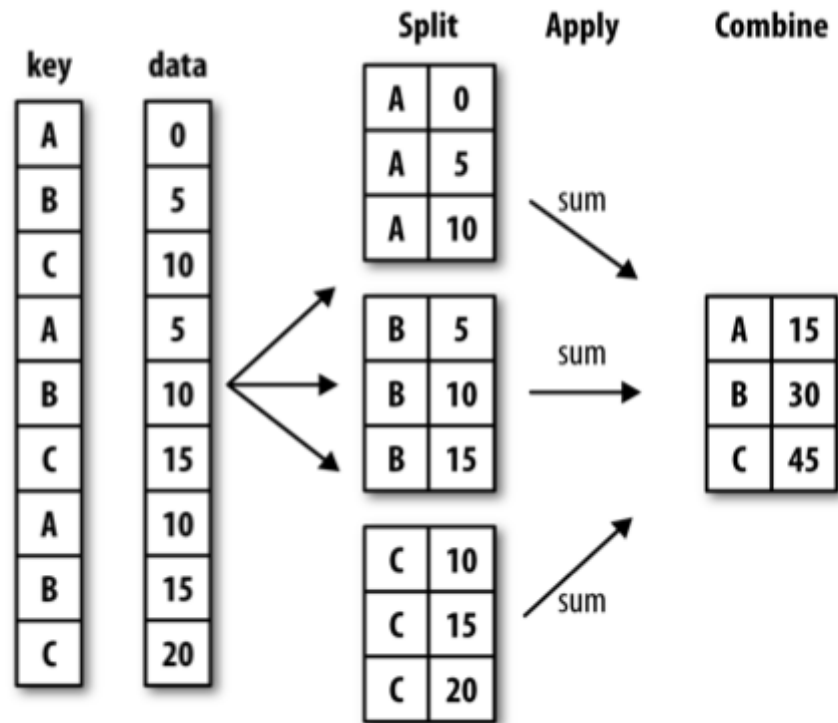
pd.options.display.max_rows = 10
```

## Recap: the groupby operation (split-apply-combine)

The "group by" concept: we want to **apply the same function on subsets of your dataframe, based on some key to split the dataframe in subsets**

This operation is also referred to as the "split-apply-combine" operation, involving the following steps:

- **Splitting** the data into groups based on some criteria
- **Applying** a function to each group independently
- **Combining** the results into a data structure



Similar to SQL GROUP BY

The example of the image in pandas syntax:

```
In [ ]: df = pd.DataFrame({'key': ['A', 'B', 'C', 'A', 'B', 'C', 'A', 'B', 'C'],
                           'data': [0, 5, 10, 5, 10, 15, 10, 15, 20]})
df
```

```
Out [ ]:   data key
0      0  A
1      5  B
2     10  C
3      5  A
```

	data	key
4	10	B
5	15	C
6	10	A
7	15	B
8	20	C

Using the filtering and reductions operations we have seen in the previous notebooks, we could do something like:

```
df[df['key'] == "A"].sum()
df[df['key'] == "B"].sum()
...
```

But pandas provides the `groupby` method to do this:

```
In [ ]: df.groupby('key').aggregate('sum') # np.sum
```

```
Out[ ]:      data
      key
      --
A      15
B      30
C      45
```

```
In [ ]: df.groupby('key').sum()
```

```
Out[ ]:      data
      key
      --
A      15
```

data	
key	
<b>B</b>	30
<b>C</b>	45

Pandas does not only let you group by a column name. In `df.groupby(grouper)` can be many things:

- Series (or string indicating a column in df)
- function (to be applied on the index)
- dict : groups by values
- levels=[], names of levels in a MultiIndex

```
In [ ]: df.groupby(lambda x: x % 2).mean()
```

```
Out[ ]: data
0      10
1      10
```

## And now applying this on some real data

These exercises are based on the [PyCon tutorial of Brandon Rhodes](#) (so all credit to him!) and the datasets he prepared for that. You can download these data from here: `titles.csv` and `cast.csv` and put them in the `/data` folder.

`cast` dataset: different roles played by actors/actresses in films

- title: title of the film
- name: name of the actor/actress
- type: actor/actress
- n: the order of the role (n=1: leading role)

```
In [ ]: cast = pd.read_csv('data/cast.csv')
cast.head()
```

Out[ ]:

	title	year	name	type	character	n
0	Suuri illusioni	1985	Homo \$	actor	Guests	22.0
1	Gangsta Rap: The Glockumentary	2007	Too \$hort	actor	Himself	NaN
2	Menace II Society	1993	Too \$hort	actor	Lew-Loc	27.0
3	Porndogs: The Adventures of Sadie	2009	Too \$hort	actor	Bosco	3.0
4	Stop Pepper Palmer	2014	Too \$hort	actor	Himself	NaN

In [ ]:

```
titles = pd.read_csv('data/titles.csv')
titles.head()
```

Out[ ]:


	title	year
0	The Rising Son	1990
1	Ashes of Kukulcan	2016
2	The Thousand Plane Raid	1969
3	Crucea de piatra	1993
4	The 86	2015

**EXERCISE:** Using groupby(), plot the number of films that have been released each decade in the history of cinema.

In [ ]:

```
titles.groupby(titles.year // 10 * 10).size().plot(kind='bar')
```


Out[ ]: &lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7f710e20d9e8&gt;



**EXERCISE:** Use `groupby()` to plot the number of "Hamlet" films made each decade.

```
In [ ]: hamlet = titles[titles['title'] == 'Hamlet']  
        hamlet.groupby(hamlet.year // 10 * 10).size().plot(kind='bar')
```

```
Out[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f710c6b4940>
```



**EXERCISE:** How many leading (n=1) roles were available to actors, and how many to actresses, in each year of the 1950s?

```
In [ ]: cast1950 = cast[cast.year // 10 == 195]
cast1950 = cast1950[cast1950.n == 1]
cast1950.groupby(['year', 'type']).size()
```

```
Out[ ]: year  type
1950  actor    604
      actress  271
1951  actor    633
      actress  272
1952  actor    591
      ...
1957  actress  284
1958  actor    694
      actress  275
1959  actor    678
```

actress 287  
dtype: int64

**EXERCISE:** List the 10 actors/actresses that have the most leading roles (n=1) since the 1990's.

```
In [ ]: cast1990 = cast[cast['year'] >= 1990]
cast1990 = cast1990[cast1990.n == 1]
cast1990.groupby('name').size().nlargest(10)
```

```
Out[ ]: name
Mohanlal      126
Mammootty     118
Akshay Kumar   87
Jayaram        76
Andy Lau       72
Ajay Devgn     69
Amitabh Bachchan 68
Eric Roberts   68
Nagarjuna Akkineni 60
Dilip          59
dtype: int64
```

**EXERCISE:** Use groupby() to determine how many roles are listed for each of The Pink Panther movies.

```
In [ ]: c = cast
c = c[c.title == 'The Pink Panther']
c = c.groupby(['year'])[['n']].max()
c
```

```
Out[ ]:      n
year
1963  15.0
2006  50.0
```

**EXERCISE:** List, in order by year, each of the films in which Frank Oz has played more than 1 role.



```
In [ ]: c = cast
c = c[c.name == 'Frank Oz']
g = c.groupby(['year', 'title']).size()
g[g > 1]
```

```
Out[ ]: year  title
1979  The Muppet Movie      8
1981  An American Werewolf in London  2
      The Great Muppet Caper      6
1982  The Dark Crystal      2
1984  The Muppets Take Manhattan      7
1985  Sesame Street Presents: Follow that Bird  3
1992  The Muppet Christmas Carol      7
1996  Muppet Treasure Island      4
1999  Muppets from Space      4
      The Adventures of Elmo in Grouchland  3
dtype: int64
```

**EXERCISE:** List each of the characters that Frank Oz has portrayed at least twice.

```
In [ ]: c = cast
c = c[c.name == 'Frank Oz']
g = c.groupby(['character']).size()
g[g > 1].sort_values()
```

```
Out[ ]: character
Grover      2
Bert        3
Cookie Monster  3
Fozzie Bear  4
Sam the Eagle  5
Yoda        5
Animal      6
Miss Piggy   6
dtype: int64
```

## Transforms

Sometimes you don't want to aggregate the groups, but transform the values in each group. This can be achieved with `transform` :

In [ ]:

df

Out[ ]:

	data	key
0	0	A
1	5	B
2	10	C
3	5	A
4	10	B
5	15	C
6	10	A
7	15	B
8	20	C

In [ ]:

`df.groupby('key').transform('mean')`

Out[ ]:

	data
0	5
1	10
2	15
3	5
4	10
5	15
6	5
7	10
8	15

```
In [ ]: def normalize(group):  
        return (group - group.mean()) / group.std()
```

```
In [ ]: df.groupby('key').transform(normalize)
```

```
Out[ ]: data  
0    -1.0  
1    -1.0  
2    -1.0  
3     0.0  
4     0.0  
5     0.0  
6     1.0  
7     1.0  
8     1.0
```

```
In [ ]: df.groupby('key').transform('sum')
```

```
Out[ ]: data  
0    15  
1    30  
2    45  
3    15  
4    30  
5    45  
6    15
```

**data**

7	30
8	45

**EXERCISE:** Add a column to the *c* \* dataframe that indicates the number of roles for the film.

```
In [ ]: cast['n_total'] = cast.groupby('title')['n'].transform('max')
cast.head()
```

```
Out[ ]:
```

	title	year	name	type	character	n	n_total
0	Suuri illusioni	1985	Homo \$	actor	Guests	22.0	22.0
1	Gangsta Rap: The Glockumentary	2007	Too \$hort	actor	Himself	NaN	NaN
2	Menace II Society	1993	Too \$hort	actor	Lew-Loc	27.0	45.0
3	Porndogs: The Adventures of Sadie	2009	Too \$hort	actor	Bosco	3.0	9.0
4	Stop Pepper Palmer	2014	Too \$hort	actor	Himself	NaN	NaN

**EXERCISE:** Calculate the ratio of leading actor and actress roles to the total number of leading roles per decade.

Tip: you can to do a groupby twice in two steps, once calculating the numbers, and then the ratios.

```
In [ ]: leading = cast[cast['n'] == 1]
sums_decade = leading.groupby([cast['year'] // 10 * 10, 'type']).size()
sums_decade
```

```
Out[ ]:
```

year	type	
1900	actor	5
	actress	1
1910	actor	2406
	actress	2753
1920	actor	4485
	...	
2000	actress	7537

```

2010  actor      17262
      actress    7384
2020  actor        3
      actress     1
dtype: int64

```

```

In [ ]: #sums_decade.groupby(level='year').transform(lambda x: x / x.sum())
        ratios_decade = sums_decade / sums_decade.groupby(level='year').transform('sum')
        ratios_decade

```

```

Out[ ]: year  type
1900  actor      0.833333
      actress    0.166667
1910  actor      0.466369
      actress    0.533631
1920  actor      0.598080
      ...
2000  actress    0.295464
2010  actor      0.700398
      actress    0.299602
2020  actor      0.750000
      actress    0.250000
dtype: float64

```

```


In [ ]: ratios_decade[:, 'actor'].plot()
        ratios_decade[:, 'actress'].plot()

```

```

Out[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f710c5b3e80>

```



## Intermezzo: string manipulations

Python strings have a lot of useful methods available to manipulate or check the content of the string:

```
In [ ]: s = 'Bradwurst'
```

```
In [ ]: s.startswith('B')
```

```
Out[ ]: True
```

In pandas, those methods (together with some additional methods) are also available for string Series through the `.str` accessor:

```
In [ ]: s = pd.Series(['Bradwurst', 'Kartoffelsalat', 'Sauerkraut'])
```

```
In [ ]: s.str.startswith('B')
```

```
Out[ ]: 0    True
        1   False
        2   False
        dtype: bool
```

For an overview of all string methods, see: <http://pandas.pydata.org/pandas-docs/stable/api.html#string-handling>

**EXERCISE:** We already plotted the number of 'Hamlet' films released each decade, but not all titles are exactly called 'Hamlet'. Give an overview of the titles that contain 'Hamlet', and that start with 'Hamlet':

```
In [ ]: hamlets = titles[titles['title'].str.contains('Hamlet')]
        hamlets['title'].value_counts()
```

```
Out[ ]: Hamlet                19
        Hamlet (II)           5
        Hamlet (III)          2
        Han, hun og Hamlet     2
        Fuck Hamlet           1
        ..
        Hamlet: Prince of Denmark 1
        Zombie Hamlet          1
        Hamlet_X               1
        Dogg's Hamlet, Cahoot's Macbeth 1
        Predstava 'Hamleta' u Mrdusi Donjoj 1
        Name: title, dtype: int64
```

```
In [ ]: hamlets = titles[titles['title'].str.match('Hamlet')]
        hamlets['title'].value_counts()
```

```
Out[ ]: Hamlet                19
        Hamlet (II)           5
        Hamlet (III)          2
        Hamlet the Vampire Slayer 1
        Hamlet's Ghost         1
        ..
        Hamlet: Prince of Denmark 1
        Hamlet (A Modern Adaptation) 1
        Hamlet_X               1
        Hamlet: The Fall of a Sparrow 1
        Hamlet in the Hamptons      1
        Name: title, dtype: int64
```

**EXERCISE:** List the 10 movie titles with the longest name.

```
In [ ]: title_longest = titles['title'].str.len().nlargest(10)
        title_longest
```

```
Out[ ]: 127048    208
        28483     196
        103597    116
        8396      114
        85693     104
        108020    104
        206303     101
        122757     99
        52929     94
        187654     92
        Name: title, dtype: int64
```

```
In [ ]: pd.options.display.max_colwidth = 210
        titles.loc[title_longest.index]
```

```
Out[ ]:
```

		title	year
127048	Night of the Day of the Dawn of the Son of the Bride of the Return of the Revenge of the Terror of the Attack of the Evil Mutant Hellbound Flesh Eating	Crawling Alien Zombified Subhumanoid Living Dead, Part 5	2011
28483	Night of the Day of the Dawn of the Son of the Bride of the Return of the Revenge of the Terror of the Attack of the Evil, Mutant, Hellbound, Flesh-Eating	Subhumanoid Zombified Living Dead, Part 3	2005
103597	Maverick and Ariel's First Ever Ever Movie Hollywood or Else... (Ang pinakamahabang title ng movie sa balat ng lupa)		2010
8396	The Fable of the Kid Who Shifted His Ideals to Golf and Finally Became a Baseball Fan and Took the Only Known Cure		1916
85693	Film d'amore e d'anarchia, ovvero 'stamattina alle 10 in via dei Fiori nella nota casa di tolleranza...'		1973
108020	Those Magnificent Men in Their Flying Machines or How I Flew from London to Paris in 25 hours 11 minutes		1965
206303	Ontologica! or a Brief Explanation of Absolutely Everything that is Known about Absolutely Everything		2012
122757	The Official Motion Pictures of the Heavyweight Boxing Contest Between Gene Tunney and Jack Dempsey		1927
52929	Something Strange: 23 Peculiar Perspectives of Metaphysical Phenomena in a Modern American Age		2012
187654	The Personal History, Adventures, Experience, & Observation of David Copperfield the Younger		1935



## Value counts

A useful shortcut to calculate the number of occurrences of certain values is `value_counts` (this is somewhat equivalent to `df.groupby(key).size()` )

For example, what are the most occurring movie titles?

```
In [ ]: titles.title.value_counts().head()
```

```
Out[ ]: Hamlet          19
        Macbeth        14
        Carmen         14
        The Three Musketeers 12
        She            11
        Name: title, dtype: int64
```

**EXERCISE:** Which years saw the most films released?


```
In [ ]: t = titles
        t.year.value_counts().head(3)
```

```
Out[ ]: 2015    13978
        2014     8209
        2013     7664
        Name: year, dtype: int64
```

**EXERCISE:** Plot the number of released films over time

```
In [ ]: titles.year.value_counts().sort_index().plot()
```


```
Out[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f710c4ddcf8>
```



**EXERCISE:** Plot the number of "Hamlet" films made each decade.

```
In [ ]: t = titles
t = t[t.title == 'Hamlet']
(t.year // 10 * 10).value_counts().sort_index().plot(kind='bar')
```

```
Out[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f710c4aeb38>
```



**EXERCISE:** What are the 11 most common character names in movie history?

```
In [ ]: cast.character.value_counts().head(11)
```

```
Out[ ]: Himself      18928
Dancer      11070
Extra       9141
Reporter    7646
Doctor      6846
...
Student     6406
Bartender   6178
Nurse       6164
Party Guest 5917
Minor Role  5880
Name: character, dtype: int64
```

**EXERCISE:** Which actors or actresses appeared in the most movies in the year 2010?


```
In [ ]: cast[cast.year == 2010].name.value_counts().head(10)
```

```
Out[ ]: Lloyd Kaufman      23
Jagathi Sreekumar      20
Suraaj Venjarammoodu   20
Chris (II) Eddy        20
Danny Trejo            17
Matt Simpson Siegel    17
Brahmanandam           15
Joe Estevez            15
Ben (II) Bishop        15
Kyle Rea               15
Name: name, dtype: int64
```

**EXERCISE:** Plot how many roles Brad Pitt has played in each year of his career.

```
In [ ]: cast[cast.name == 'Brad Pitt'].year.value_counts().sort_index().plot()
```

```
Out[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f710c490550>
```



**EXERCISE:** What are the 10 most film titles roles that start with the word "The Life"?

```
In [ ]: c = cast
c[c.title.str.startswith('The Life')].title.value_counts().head(10)
```

```
Out[ ]: The Life of David Gale          137
The Life Aquatic with Steve Zissou    78
The Life Before Her Eyes              74
The Life of Riley                     73
The Life and Death of Peter Sellers   65
The Life and Death of Colonel Blimp   58
The Life and Hard Times of Guy Terrifico 53
The Life and Times of Judge Roy Bean  50
The Life of Emile Zola                 46
The Life of the Party                 45
Name: title, dtype: int64
```

**EXERCISE:** How many leading (n=1) roles were available to actors, and how many to actresses, in the 1950s? And in 2000s?

```
In [ ]: c = cast
c = c[c.year // 10 == 195]
c = c[c.n == 1]
c.type.value_counts()
```

```
Out[ ]: actor      6388
actress    2813
Name: type, dtype: int64
```

```
In [ ]: c = cast
c = c[c.year // 10 == 200]
c = c[c.n == 1]
c.type.value_counts()
```

```
Out[ ]: actor      17972
actress    7537
Name: type, dtype: int64
```

In [ ]: